



INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 08
SQL

Week 4



- SQL – or SEQUEL is a language used in relational databases
- **DBMS support CRUD**
 - Create, Read, Update, Delete commands
- SQL supports CRUD
 - Create, Select, Update, Delete commands
- Other info
 - You can see the 2011 standard of SQL at
 - http://www.jtc1sc32.org/doc/N2151-2200/32N2153T-text_for_ballot-FDIS_9075-1.pdf
 - Wikipedia has several sections on SQL (good for generic syntax)
 - http://en.wikipedia.org/wiki/Category:SQL_keywords



- Provides the following capabilities:
 - Data Definition Language (DDL)
 - To define and set up the database
 - CREATE, ALTER, DROP
 - Data Manipulation Language (DML)
 - To maintain and use the database
 - SELECT, INSERT, DELETE, UPDATE
 - Data Control Language (DCL)
 - To control access to the database
 - GRANT, REVOKE
 - Other Commands
 - Administer the database
 - Transaction Control

- In **Implementation** of the database
 - Take the tables we design in physical design
 - Implement these tables in the database using create commands
- In **Use** of the database
 - Use Select commands to read the data from the tables, link the tables together etc
 - Use alter, drop commands to update the database
 - Use insert, update, delete commands to change data in the database

1.

```
CREATE TABLE BankHQ (
    BankHQID INT(4) AUTO_INCREMENT,
    HQAddress VARCHAR(300) NOT NULL,
    OtherHQDetails VARCHAR(500),
    PRIMARY KEY (BankHQID)
)
```

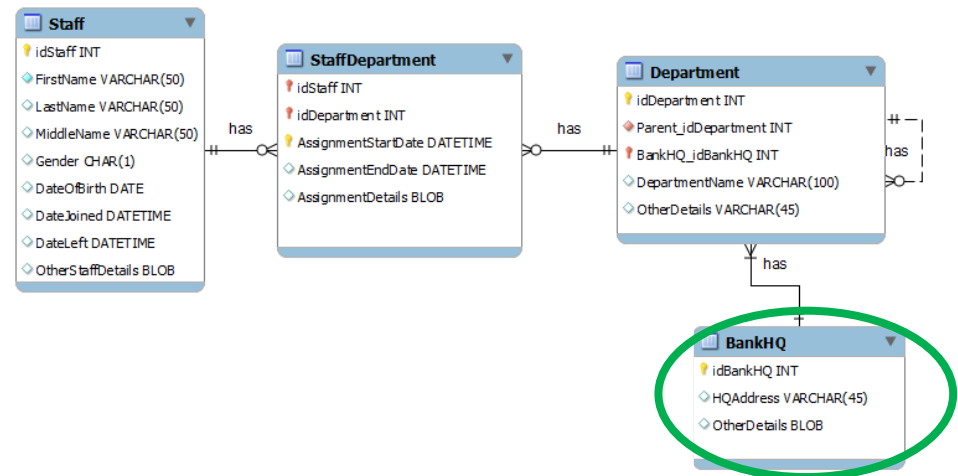
2.

```
INSERT INTO BankHQ VALUES
(DEFAULT, "23 Charles St Peterson North 2022", 'Main Branch');
INSERT INTO BankHQ VALUES
(DEFAULT, "213 Jones Rd Parkville North 2122", 'Sub Branch');
```








3.

1 • `select * from BANKHQ`

BankHQID	HQAddress	OtherHQDetails
1	23 Charles St Peterson North 2022	Main Branch
2	213 Jones Rd Parkville North 2122	Sub Branch

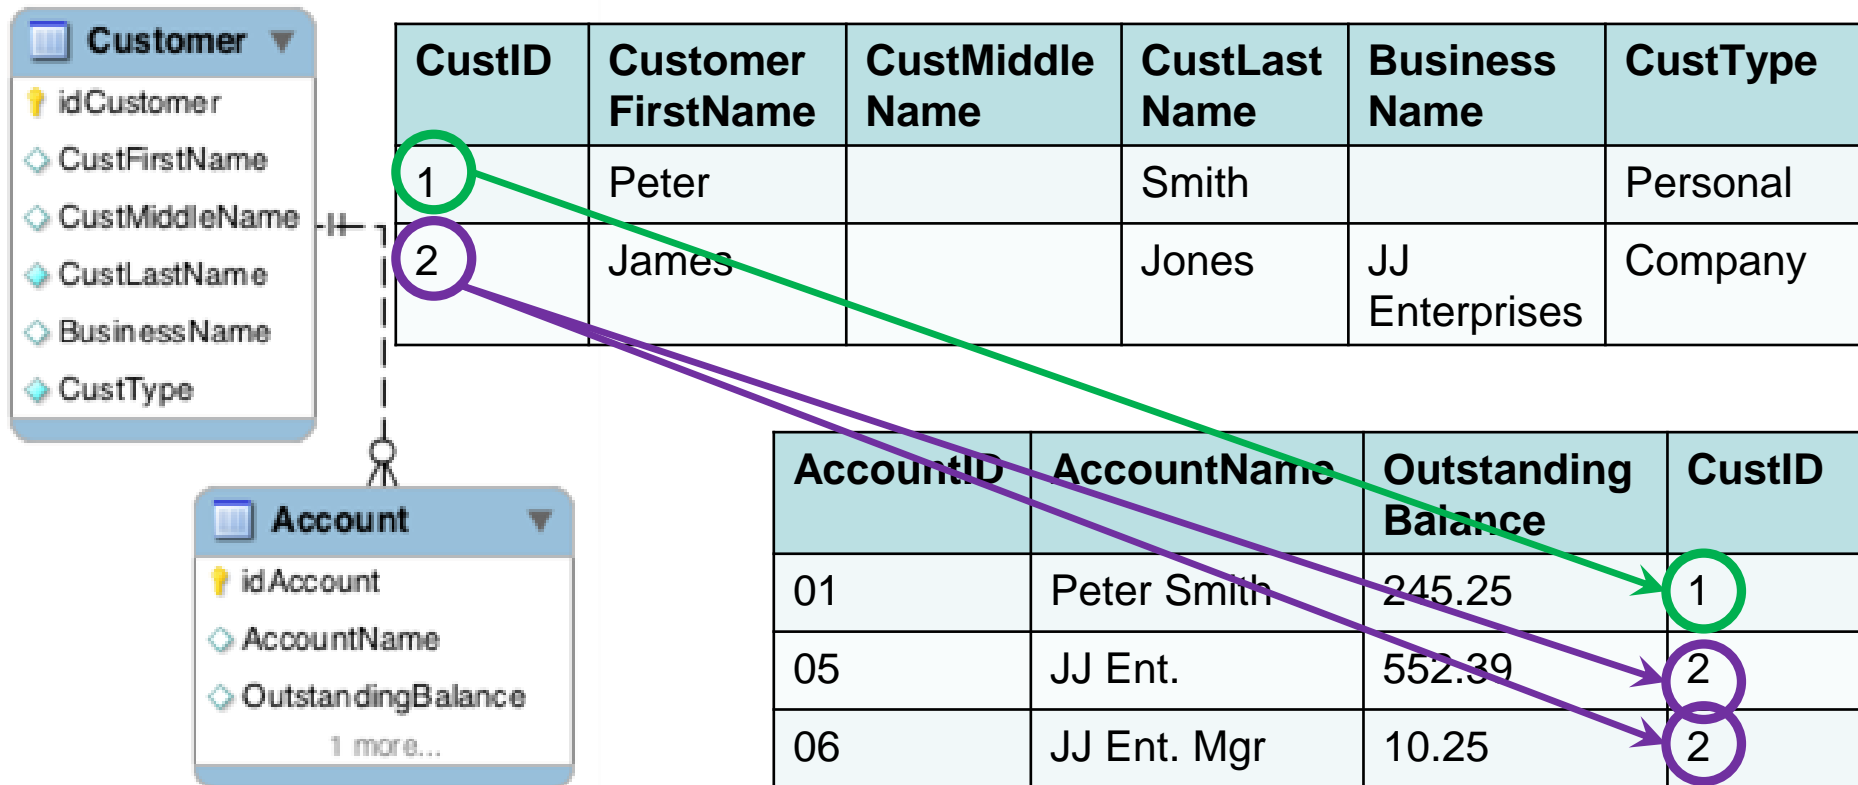


Create Table: Review

Customer	
	CustomerID
	CustFirstName
	CustMiddleName
	CustLastName
	BussinessName
	CustType
	CustAddress(Line1,Line2,Suburb,Postcode,Country...)

```
CREATE TABLE Customer (
    CustomerID          smallint          auto_increment,
    CustFirstName       varchar(100),
    CustMiddleName      varchar(100),
    CustLastName        varchar(100)      NOT NULL,
    BusinessName        varchar(200),
    CustType            enum('Personal','Company') NOT NULL,
    PRIMARY KEY (CustomerID)
);
```

- We looked at Customer
 - A customer can have a number of Accounts
 - The tables get linked through a foreign key





```
CREATE TABLE Account (  
    AccountID          smallint          auto_increment,  
    AccountName        varchar(100)      NOT NULL,  
    OutstandingBalance DECIMAL(10,2)     NOT NULL,  
    CustomerID         smallint          NOT NULL,  
    PRIMARY KEY (AccountID),  
    FOREIGN KEY (CustomerID) REFERENCES Customer(CustomerID)  
        ON DELETE RESTRICT  
        ON UPDATE CASCADE  
);
```




1) **INSERT INTO** Customer
(CustFirstName, CustLastName, CustType)
VALUES ("Peter", "Smith", 'Personal');

Specifies which columns
will be entered

2) **INSERT INTO** Customer
VALUES (DEFAULT, "James", NULL, "Jones",
"JJ Enterprises", 'Company');

No column specification means
ALL columns need to be entered

INSERT INTO Customer
VALUES (DEFAULT, "", NULL, "Smythe",
"", 'Company');

Customer

CustID	CustomerFirst Name	CustMiddle Name	CustLastName	BusinessName	CustType
1	Peter	NULL	Smith	NULL	Personal
2	James	NULL	Jones	JJ Enterprises	Company
3		NULL	Smythe		Company

What does **NULL** mean?

Null Island: The Busiest Place That Doesn't Exist:
<https://www.youtube.com/watch?v=bjvlp1-1w84>
by the channel MinuteEarth

- Select statement allows us to query table(s)
 - * (star): Allows us to obtain *all* columns from a table

All columns



```
1 • select * from Customer ;
```

Query 4 Result					
Fetches 3 records. Duration: 0.015 sec, f					
CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	Cust Type
1	Peter	NULL	Smith	NULL	Personal
2	James	NULL	Jones	JJ Enterprises	Company
3		NULL	Smythe		Company

- A cut down version of the SELECT statement – MySQL
- **SELECT** [ALL | DISTINCT] *select_expr* [, *select_expr* ...]
 - List the columns (and expressions) that are returned from the query
- **[FROM *table_references*]**
 - Indicate the table(s) or view(s) from where the data is obtained
- **[WHERE *where_condition*]**
 - Indicate the conditions on whether a particular row will be in the result
- **[GROUP BY {*col_name* | *expr* } [ASC | DESC], ...]**
 - Indicate categorisation of results
- **[HAVING *where_condition*]**
 - Indicate the conditions under which a particular category (group) is included in the result
- **[ORDER BY {*col_name* | *expr* | *position*} [ASC | DESC], ...]**
 - Sort the result based on the criteria
- **[LIMIT {[*offset*,] *row_count* | *row_count* OFFSET *offset*}]**
 - Limit which rows are returned by their return order (ie 5 rows, 5 rows from row 2)

Order is important! E.g. Limit cannot go before Group By or Having



Select Examples

SELECT * FROM Customer;
= Give me all information you have about customers

Customer
CustomerID INT
CustomerFirstName VARCHAR(100)
CustomerMiddleName VARCHAR(100)
CustomerLastName VARCHAR(100)
BussinessName VARCHAR(100)
CustomerType CHAR(1)
5 more...

SQL

SELECT * FROM Customer;					
CustomerID CustFirstName CustMiddleName CustLastName BusinessName Cust Type					
1	Peter	NULL	Smith	NULL	Personal
2	James	NULL	Jones	JJ Enterprises	Company
3	Akin	NULL	Smithies	Bay Wart	Company
4	Julie	Anne	Smythe	Konks	Company
5	Jen	NULL	Smart	BRU	Company
6	Lim	NULL	Lam	NULL	Personal
7	Kim	NULL	Unila	Saps	Company
8	James	Jay	Jones	JJ's	Company
9	Keith	NULL	Samson	NULL	Personal
NULL	NULL	NULL	NULL	NULL	NULL

RESULT

Customer	
CustomerID	INT
CustomerFirstName	VARCHAR(100)
CustomerMiddleName	VARCHAR(100)
CustomerLastName	VARCHAR(100)
BusinessName	VARCHAR(100)
CustomerType	CHAR(1)
5 more...	

In Relational Algebra:

$$\pi_{CustLastName}(Customer)$$

In SQL:

SELECT CustLastName
FROM Customer;

NOTE: MySQL doesn't discard duplicates.
To remove them use DISTINCT in front of
the projection list.

SQL

SELECT CustLastName FROM Customer;	
Result	
CustLastName	
Smith	
Jones	
Smithies	
Smythe	
Smart	
Lam	
Unila	
Jones	
Samson	

In Relational Algebra:

$$\sigma_{cond1 \wedge cond2 \vee cond3}^{(Rel)}$$

In SQL:

WHERE cond1 AND cond2
OR cond3

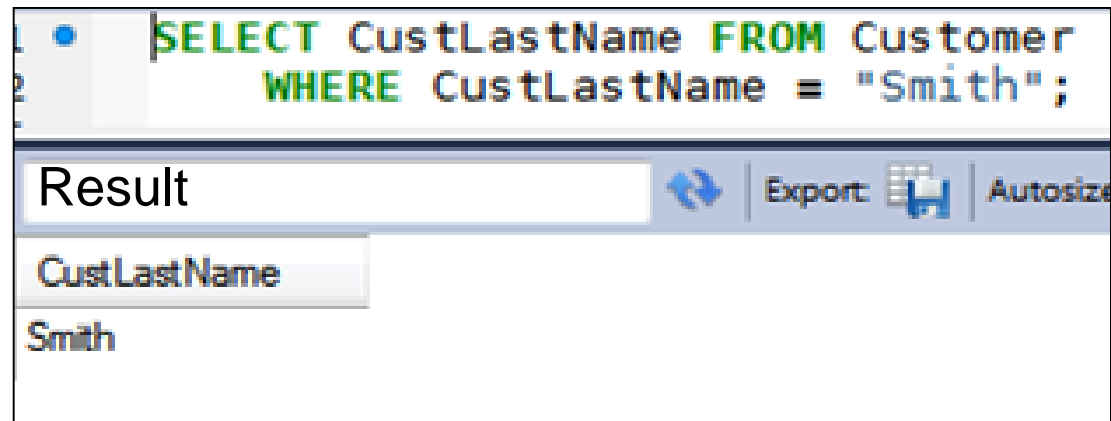
In Relational Algebra:

$$\pi_{CustLastName}(\sigma_{CustLastName="Smith"}(Customer))$$

In SQL:

SELECT CustLastName
FROM Customer
WHERE CustLastName = "Smith";

SQL





- In addition to arithmetic expressions, string conditions are specified with the LIKE clause

LIKE “REG_EXP”

% Represents zero, one, or multiple characters

_ Represents a single character

Examples:

WHERE CustomerName LIKE 'a%'	Finds any values that start with "a"
WHERE CustomerName LIKE '%a'	Finds any values that end with "a"
WHERE CustomerName LIKE '%or%'	Finds any values that have "or" in any position
WHERE CustomerName LIKE '_r%'	Finds any values that have "r" in the second position
WHERE CustomerName LIKE 'a_%_%'	Finds any values that start with "a" and are at least 3 characters in length
WHERE ContactName LIKE 'a%o'	Finds any values that start with "a" and end with "o"

SQL:

```
SELECT CustLastName FROM Customer  
WHERE CustLastName LIKE "Sm%";
```

CustLastName
Smith
Smithies
Smythe
Smart

Column renaming

We can rename the column name of the output by using the AS clause

```
SELECT CustType, Count(CustomerID)
FROM Customer
GROUP BY CustType;
```

Cust Type	Count(CustomerID)
Personal	3
Company	6

```
SELECT CustType, Count(CustomerID) AS Count
FROM Customer
GROUP BY CustType;
```

Cust Type	Count
Personal	3
Company	6



Aggregate functions operate on the (sub)set of values in a column of a relation (table) and return a single value

- AVG()
 - Average value
- MIN()
 - Minimum value
- MAX()
 - Maximum value
- COUNT()
 - Number of values
- SUM()
 - Sum of values

- Plus others

- <http://dev.mysql.com/doc/refman/5.5/en/group-by-functions.html>

- All of these except for COUNT(*) ignore null values and return null if all values are null. COUNT(*) counts the number of records.



COUNT() - returns the number of records
AVG() - average of the values

Examples:

SELECT COUNT(CustomerID)
FROM Customer; = How many customers do we have
(cardinality)

SELECT AVG(OutstandingBalance)
FROM Account; = What is the average balance of
ALL ACCOUNTS

SELECT AVG(OutstandingBalance)
FROM Account
WHERE CustomerID= 1; = What is the average balance of
Accounts of Customer 1

SELECT AVG(OutstandingBalance)
FROM Account
GROUP BY CustomerID; = What is the average balance
PER CUSTOMER



- **Group by** groups all records together over a set of attributes
- Frequently used with aggregate functions

- **Example:**

*What is the average balance **PER CUSTOMER***

```
SELECT AVG(OutstandingBalance)
```

```
FROM Account
```

```
GROUP BY CustomerID;
```

Returns one record per each customer



- The HAVING clause was added to SQL because the WHERE keyword **cannot be used** with **aggregate** functions

```
SELECT column_name(s)
FROM table_name
WHERE condition
GROUP BY column_name(s)
HAVING condition
ORDER BY column_name(s);
```

- Example:**

List the number of customers of each country, but ONLY include countries with more than 5 customers

```
SELECT COUNT(CustomerID), CountryName
FROM Customers
GROUP BY CountryName
HAVING COUNT(CustomerID) > 5;
```

Condition over the aggregate

- Orders records by particular column(s)

ORDER BY XXX ASC/DESC (ASC is default)

SQL

```
SELECT CustLastName, CustType
FROM Customer
ORDER BY CustLastName;
```

RESULT

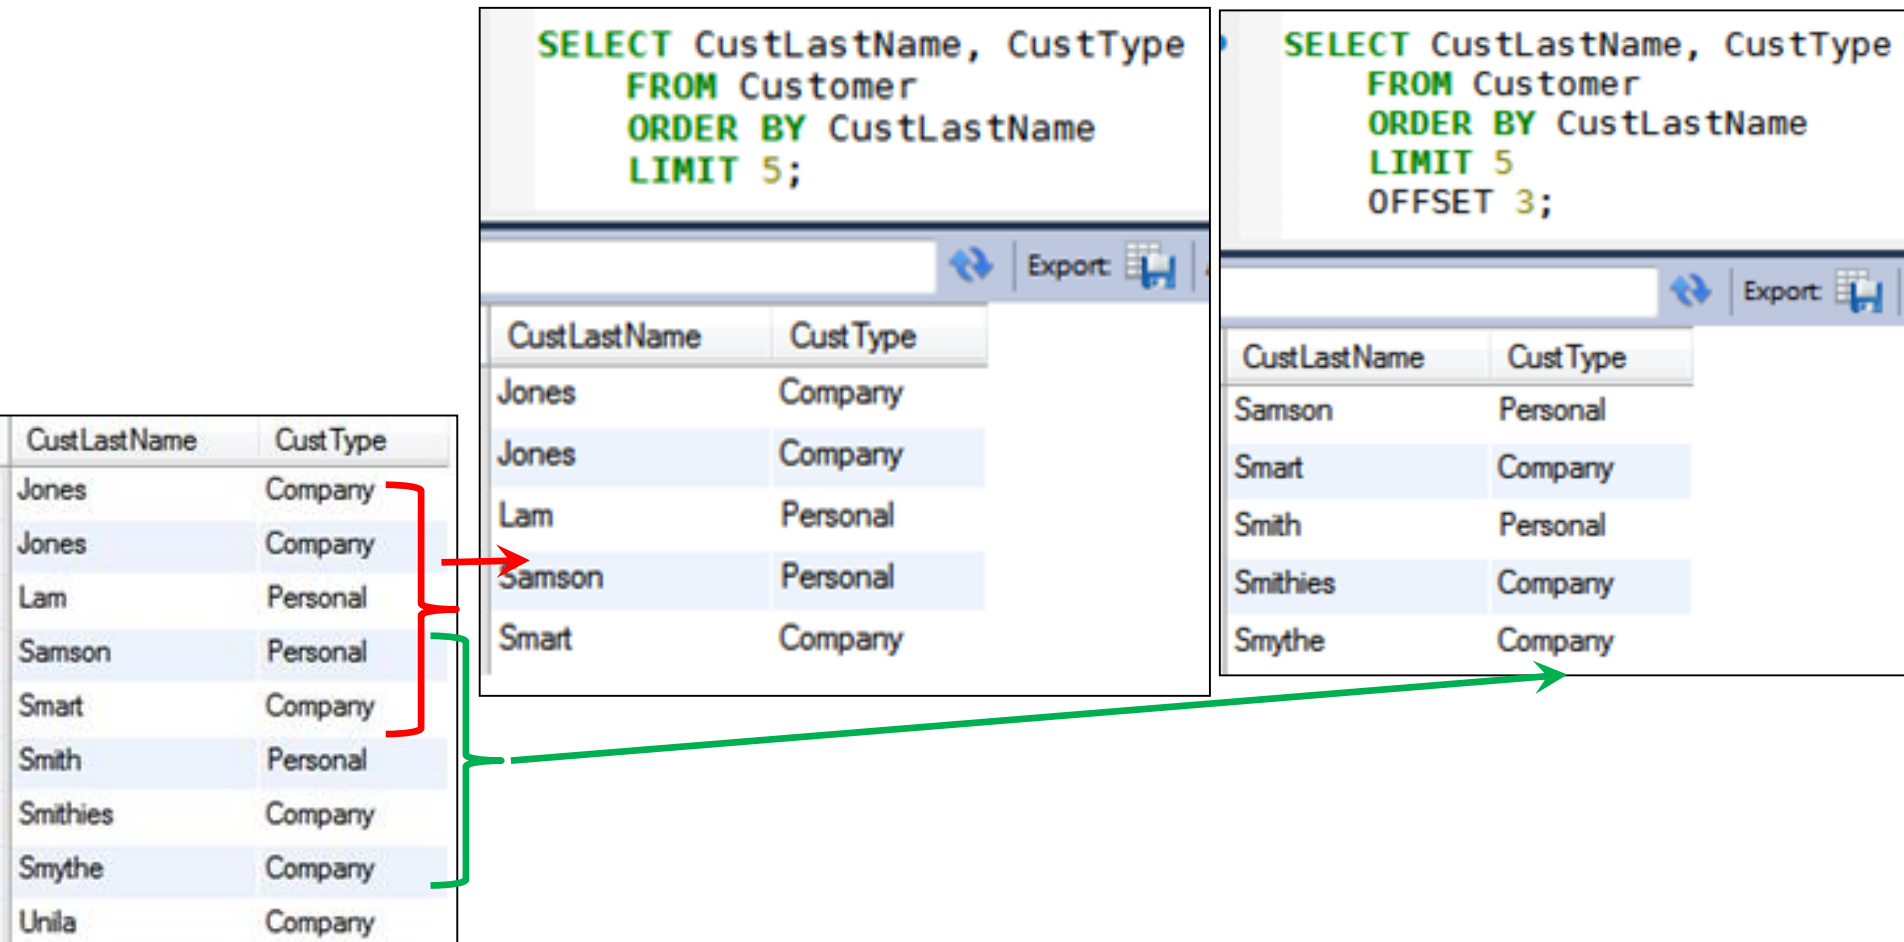
CustLastName	Cust Type
Jones	Company
Jones	Company
Lam	Personal
Samson	Personal
Smart	Company
Smith	Personal
Smithies	Company
Smythe	Company
Unila	Company

```
SELECT CustLastName, CustType
FROM Customer
ORDER BY CustLastName DESC;
```

CustLastName	Cust Type
Unila	Company
Smythe	Company
Smithies	Company
Smith	Personal
Smart	Company
Samson	Personal
Lam	Personal
Jones	Company
Jones	Company

Limit and Offset

- LIMIT number - limits the output size
- OFFSET number - skips first 'number' records



Full Customer List:

CustLastName	Cust Type
Jones	Company
Jones	Company
Lam	Personal
Samson	Personal
Smart	Company
Smith	Personal
Smithies	Company
Smythe	Company
Unila	Company

Query 1: LIMIT 5

```
SELECT CustLastName, CustType
FROM Customer
ORDER BY CustLastName
LIMIT 5;
```

CustLastName	Cust Type
Jones	Company
Jones	Company
Lam	Personal
Samson	Personal
Smart	Company

Query 2: LIMIT 5, OFFSET 3

```
SELECT CustLastName, CustType
FROM Customer
ORDER BY CustLastName
LIMIT 5
OFFSET 3;
```

CustLastName	Cust Type
Samson	Personal
Smart	Company
Smith	Personal
Smithies	Company
Smythe	Company



- `SELECT * FROM Rel1, Rel2;` - this is a **cross product**

```
SELECT * FROM Customer, Account;
```

Export: Autosize:									
CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	1	Peter Smith	245.25	1
3	Akin	NULL	Smithies	Bay Wart	Company	1	Peter Smith	245.25	1
1	Peter	NULL	Smith	NULL	Personal	2	JJ Ent.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ Ent.	552.39	2
3	Akin	NULL	Smithies	Bay Wart	Company	2	JJ Ent.	552.39	2
1	Peter	NULL	Smith	NULL	Personal	3	JJ Ent. Mgr	10.25	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ Ent. Mgr	10.25	2
3	Akin	NULL	Smithies	Bay Wart	Company	3	JJ Ent. Mgr	10.25	2

Not quite useful...

Typically we would like to find:

For every record in the Customer table list every record in the Account table



- **Inner/Equi join:**
 - Joins the tables over keys

```
SELECT * FROM Customer INNER JOIN Account  
ON Customer.CustomerID = Account.CustomerID; CONDITION
```

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ Ent.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ Ent. Mgr	10.25	2

- **Natural Join:**
 - Joins the tables over keys. The condition does not have to be specified (natural join does it automatically), but key attributes have to have the *same name*.

```
SELECT * FROM Customer NATURAL JOIN Account;
```

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	CustType	AccountID	AccountName	OutstandingBalance
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ Ent.	552.39
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ Ent. Mgr	10.25



- **Outer join:**
 - Joins the tables over keys
 - Can be *left* or *right* (see difference below)
 - Includes records that **don't match** the join from the other table

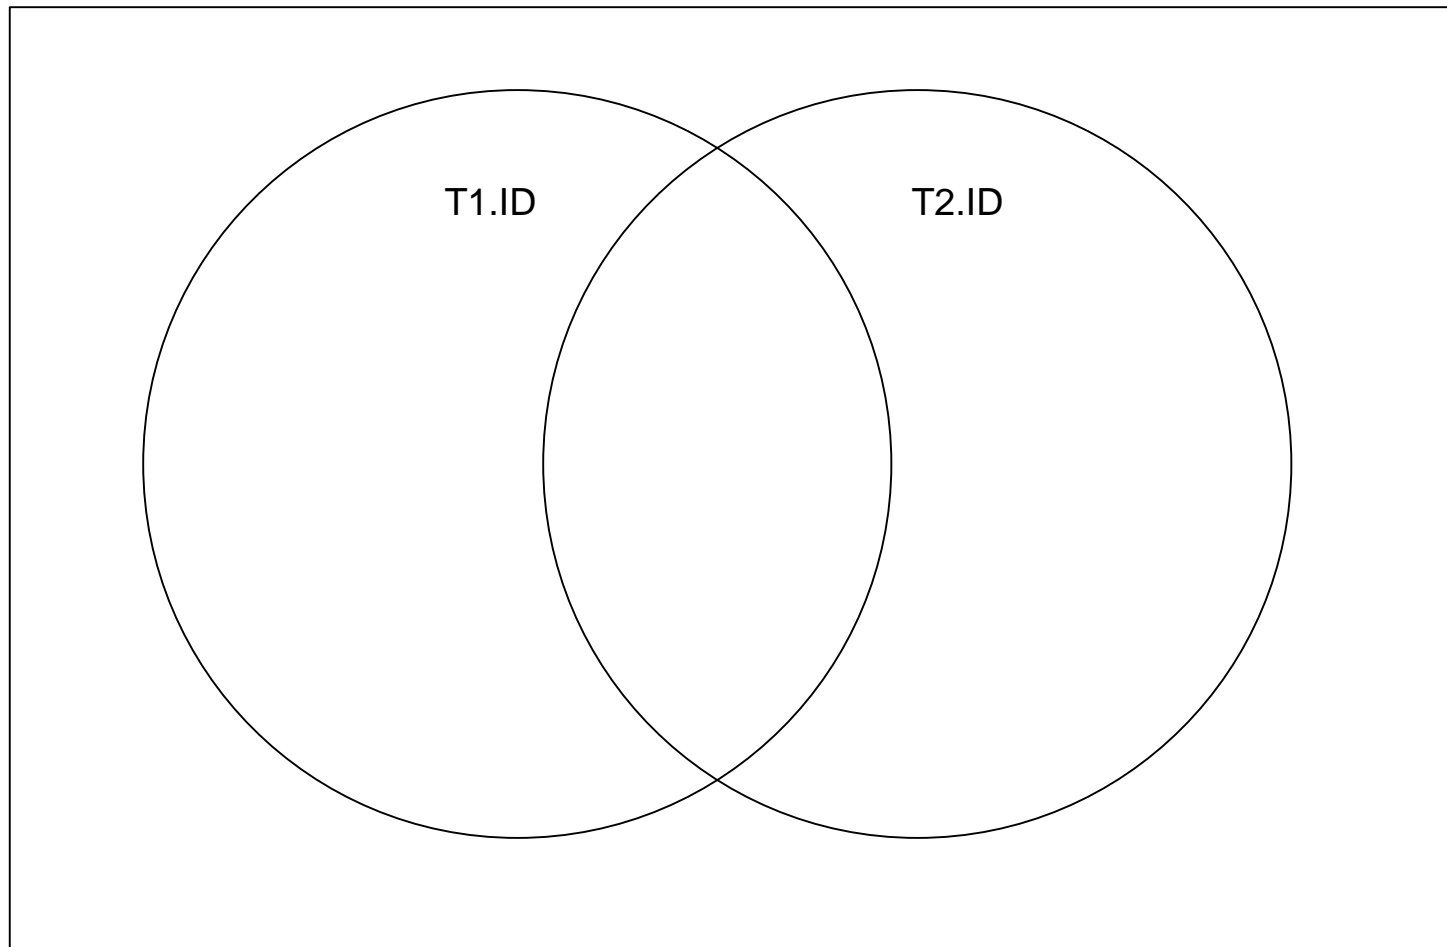
```
SELECT * FROM Customer LEFT OUTER JOIN Account
ON Customer.CustomerID = Account.CustomerID;
```

CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	Cust Type	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ ENT.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ ENT. Mgr	10.25	2
3	Akin	NULL	Smithies	Bay Wart	Company	NULL	NULL	NULL	NULL

```
SELECT * FROM Customer RIGHT OUTER JOIN Account
ON Customer.CustomerID = Account.CustomerID;
```

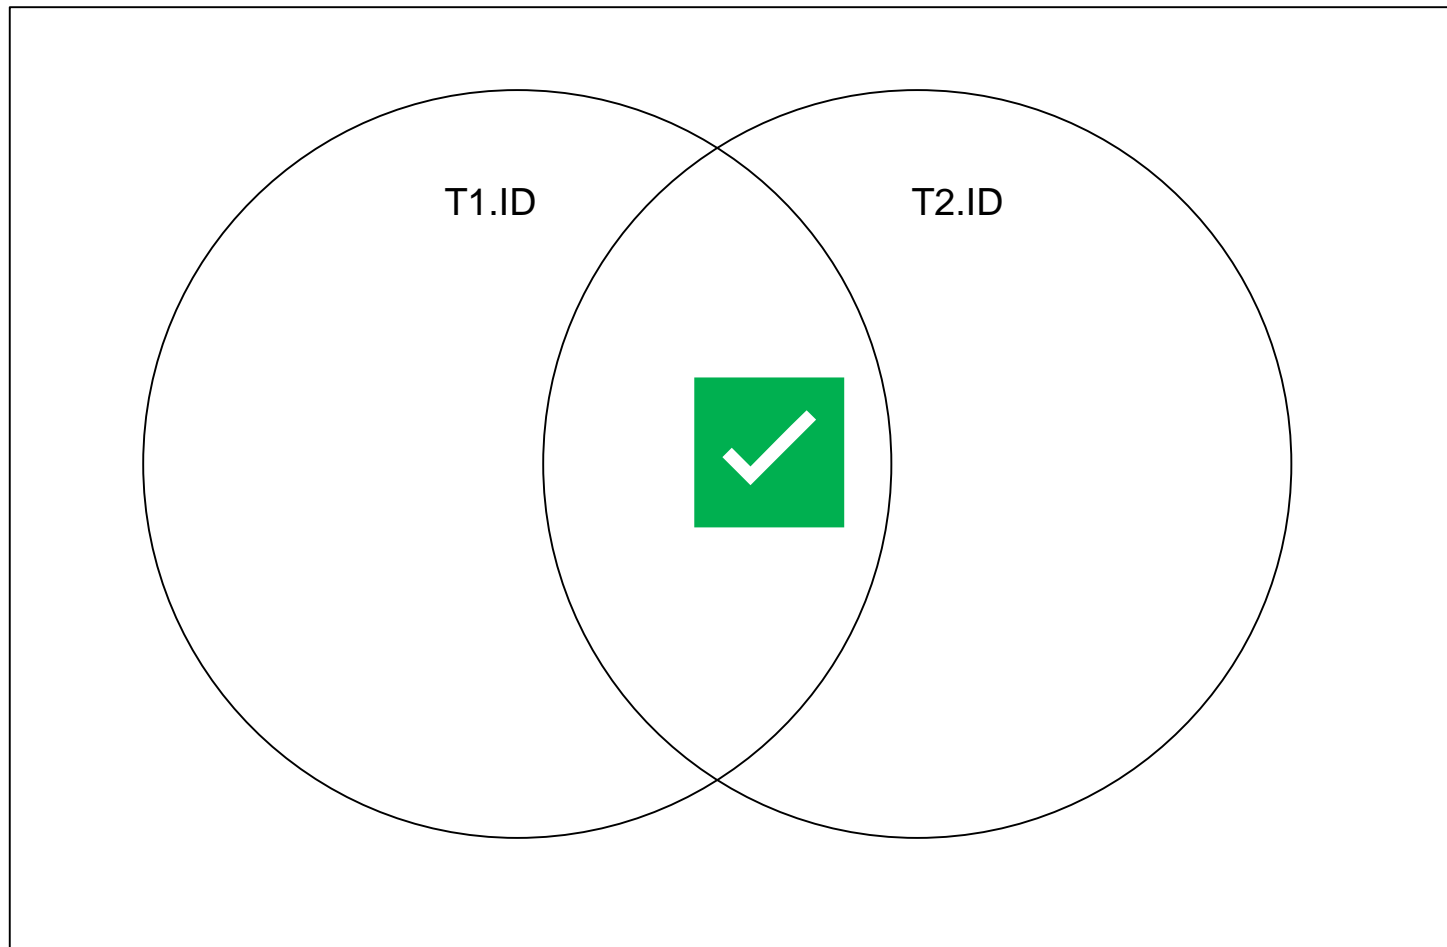
CustomerID	CustFirstName	CustMiddleName	CustLastName	BusinessName	Cust Type	AccountID	AccountName	OutstandingBalance	CustomerID
1	Peter	NULL	Smith	NULL	Personal	1	Peter Smith	245.25	1
2	James	NULL	Jones	JJ Enterprises	Company	2	JJ ENT.	552.39	2
2	James	NULL	Jones	JJ Enterprises	Company	3	JJ ENT. Mgr	10.25	2

JOINS depicted as Venn Diagrams

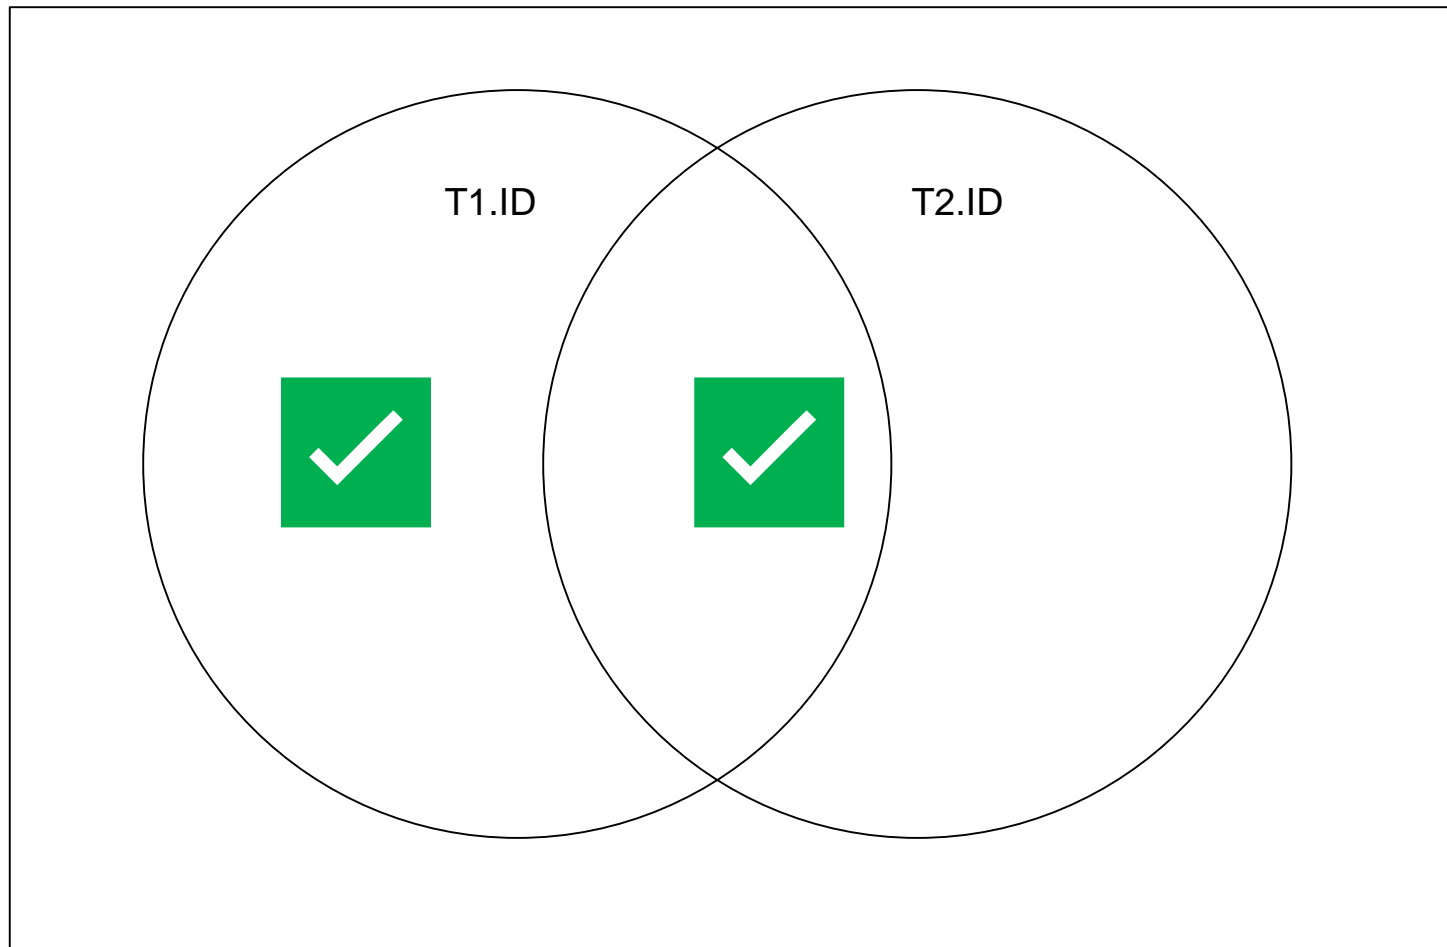


JOINS depicted as Venn Diagrams

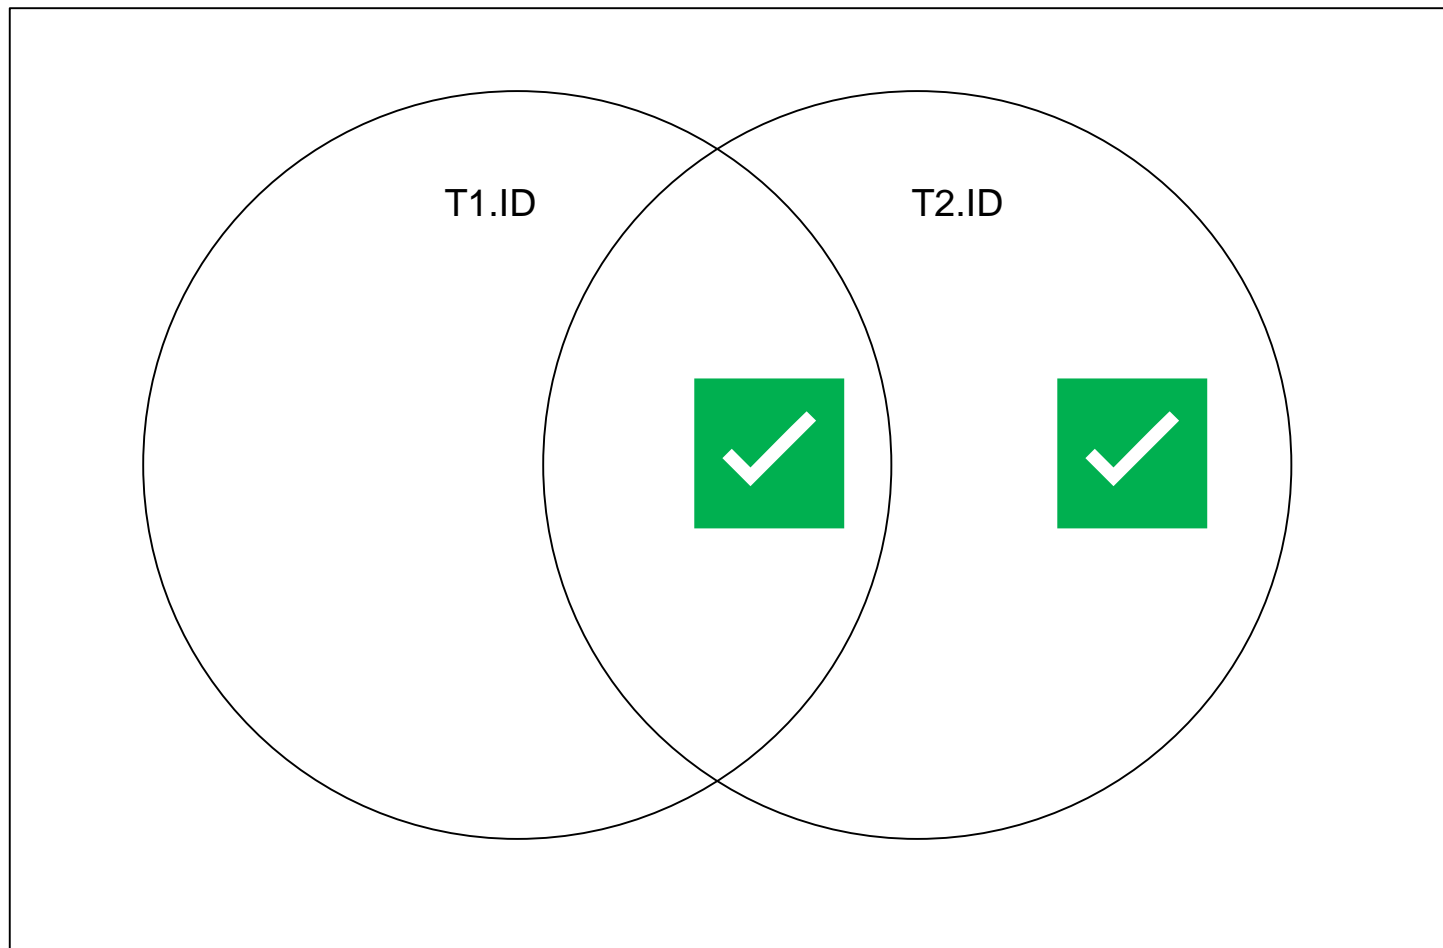
- T1 INNER JOIN T2 ON T1.ID = T2.ID
- T1 NATURAL JOIN T2



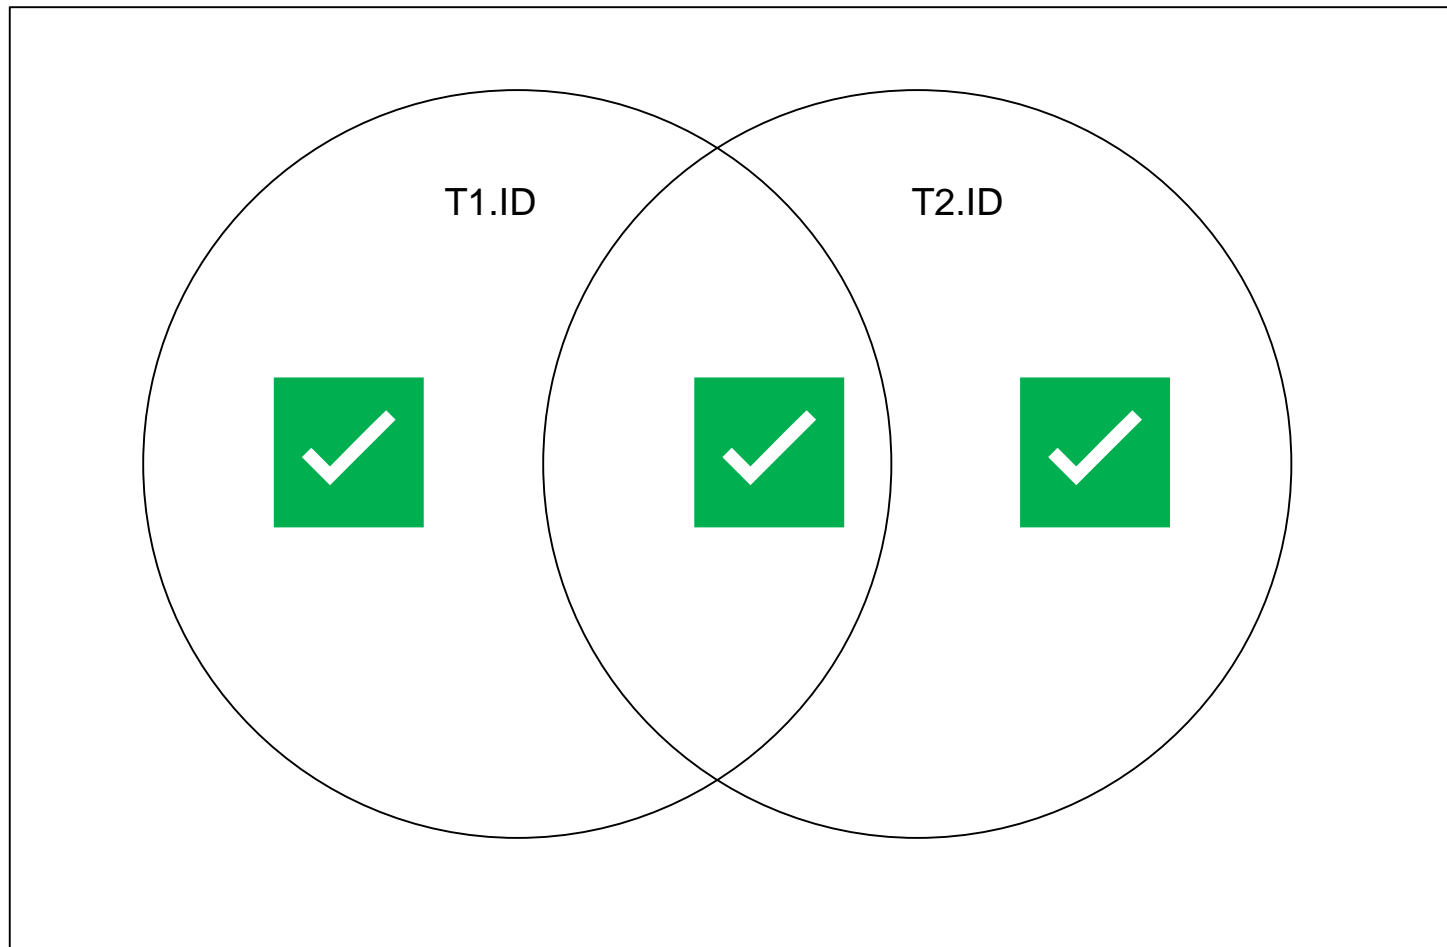
- T1 LEFT OUTER JOIN T2 ON T1.ID = T2.ID



- T1 RIGHT OUTER JOIN T2 ON T1.ID = T2.ID



- T1 FULL OUTER JOIN T2 ON T1.ID = T2.ID





- You need to know how to write SQL
 - DDL
 - DML



- SQL Summary
 - Overview of concepts, more examples