



INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 9
SQL Summary

Week 5



- Extending your knowledge
 - DML
 - Comparison & Logic Operators
 - Set Operations
 - Subquery
 - Multiple record INSERTs, INSERT from a table
 - UPDATE, DELETE, REPLACE
 - Views
 - DDL
 - ALTER and DROP, TRUNCATE, RENAME
 - DCL
- How to think about SQL
 - Problem Solving



MELBOURNE

- SQL keywords are case insensitive
 - We try to CAPITALISE them to make them clear
- Table names are Operating System Sensitive
 - If case sensitivity exists in the operating system, then the table names are case sensitive! (i.e. Mac, Linux)
 - Account <> ACCOUNT
- Field names are case insensitive
 - ACCOUNTID == AccountID == AcCoUnTID
- You can do maths in SQL...
 - SELECT 1*1+1/1-1;

MELBOURNE

- Comparison:

Operator	Description
=	Equal to
<	Less than
>	Greater than
<=	Less than or equal to
>=	Greater than or equal to
<> OR !=	Not equal to (depends on DBMS which is used)

- Logic:

- AND, NOT, OR
- **Example:** SELECT * FROM Furniture WHERE ((Type="Chair" AND Colour = "Black") OR (Type = "Lamp" AND Colour = "Black"))



- UNION
 - Shows all rows returned from the queries (or tables)
- INTERSECT
 - Shows only rows that are common in the queries (or the tables)
- [UNION/INTERSECT] ALL
 - If you want duplicate rows shown in the results you need to use the ALL keyword.. UNION ALL etc.
- In MySQL only UNION and UNION ALL are supported

UNION Example

```
SELECT Employee.Name, EmployeeType
FROM Employee INNER JOIN Hourly
ON Employee.ID = Hourly.ID

UNION

SELECT Employee.Name, EmployeeType
FROM Employee INNER JOIN Salaried
ON Employee.ID = Salaried.ID;
```

The screenshot shows a database query result window titled "Lecture7.sql Result". The window displays a table with two columns: "Name" and "Employee Type". The table contains four rows of data:

Name	Employee Type
Alice	H
Alan	H
Sean	S
Linda	S

The status bar at the bottom indicates "Fetched 4 records. Duration: 0.000s".

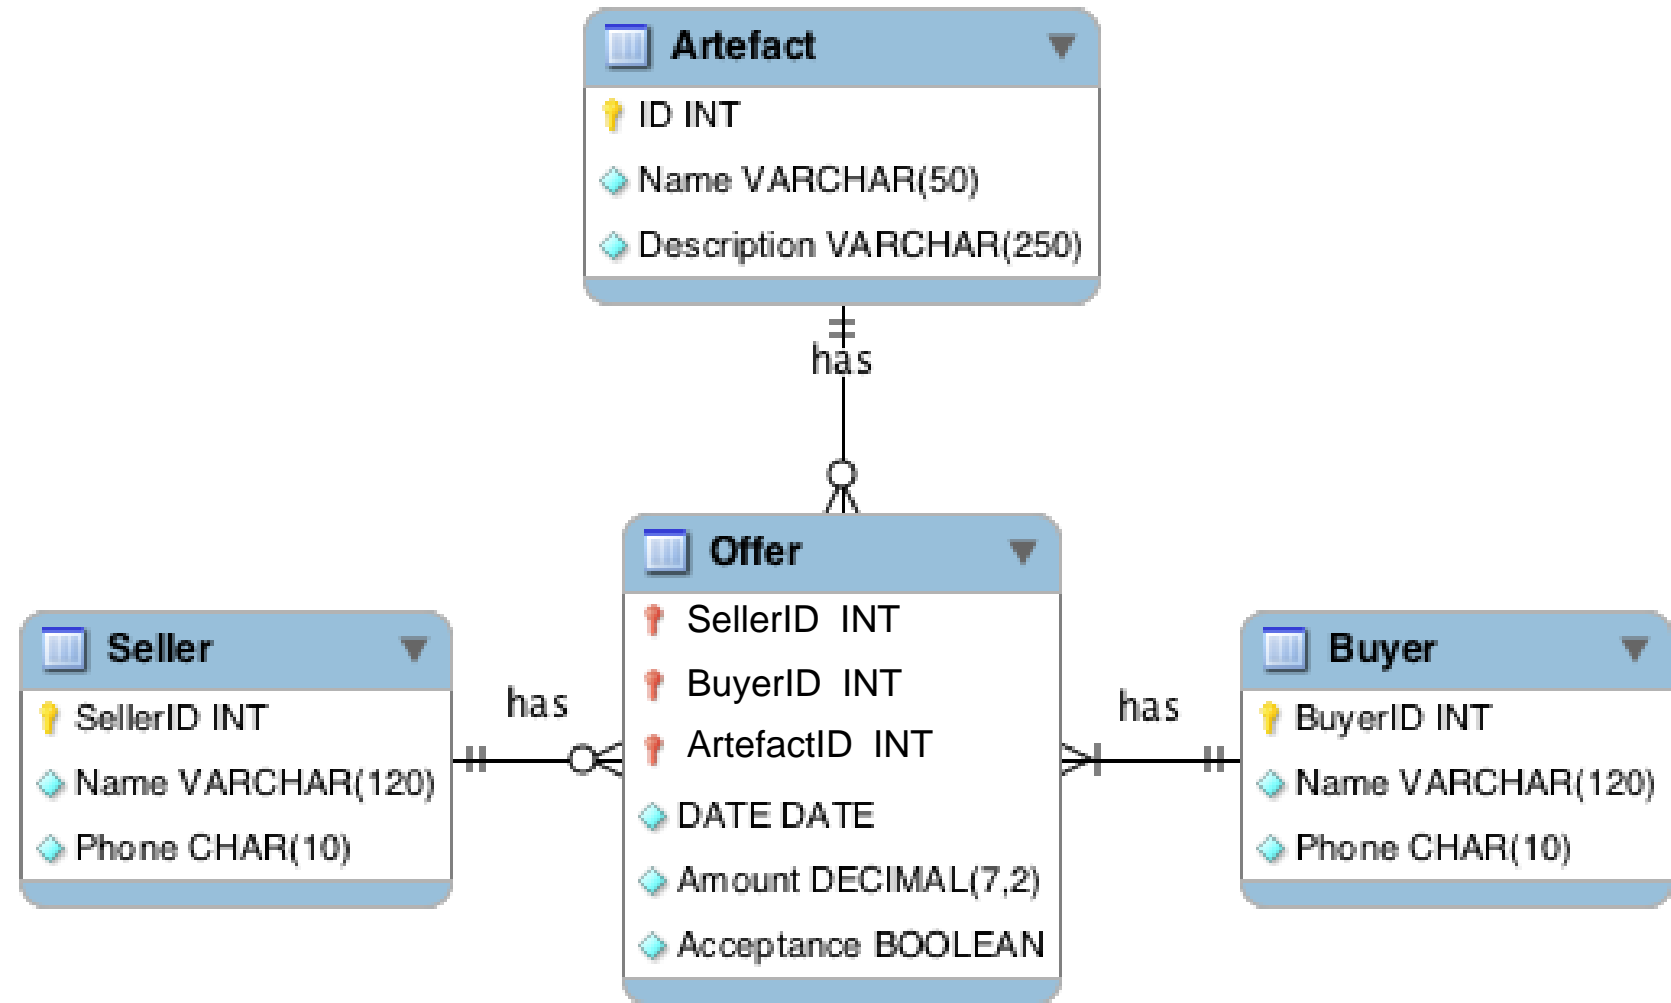


MELBOURNE

- SQL provides the ability to *nest* subqueries
- A nested query is simply another select query you write to produce a table set
 - Remember that all select queries return a table set of data
- A common use of subqueries is to perform set tests
 - Set membership, set comparisons



- IN / NOT IN
 - Used to test whether the attribute is IN/NOT IN the subquery list
- ANY
 - True if any value returned meets the condition
- ALL
 - True if all values returned meet the condition
- EXISTS
 - True if the subquery returns one or more records
- For more info:
- https://www.w3schools.com/sql/sql_any_all.asp
- https://www.w3schools.com/sql/sql_exists.asp
- General help with SQL: <https://www.w3schools.com/sql/> **(great tutorial!)**



Seller

SellerID	Name	Phone
1	Abby	0233232232
2	Ben	0311111111
3	Carl	0333333333

Offer

SellerID	ArtefactID	BuyerID	Date	Amount	Acceptance
1	1	1	2012-06-20	81223.23	N
1	1	2	2012-06-20	82223.23	N
2	2	1	2012-06-20	19.95	N
2	2	2	2012-06-20	23.00	N

Artefact

ID	Name	Description
1	Vase	Old Vase
2	Knife	Old Knife
3	Pot	Old Pot

Buyer

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444
3	Oleg	0555555555



Example: Subquery

List the BuyerID, Name and Phone number for all bidders on artefact 1

```
SELECT * FROM Buyer
WHERE BuyerID IN
(SELECT BuyerID FROM Offer WHERE ArtefactID = 1)
```

Offer

SellerID	ArtefactID	BuyerID	Date	Amount	Acceptance
1	1	1	2012-06-20	81223.23	N
1	1	2	2012-06-20	82223.23	N
2	2	1	2012-06-20	19.95	N
2	2	2	2012-06-20	23.00	N

Buyer

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444
3	Oleg	0555555555

Result

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444

Which Artefacts don't have offers made on them?

```
SELECT * FROM Artefact
WHERE ID NOT IN
(SELECT ArtefactID FROM Offer);
```

Offer

SellerID	ArtefactID	BuyerID	Date	Amount	Acceptance
1	1	1	2012-06-20	81223.23	N
1	1	2	2012-06-20	82223.23	N
2	2	1	2012-06-20	19.95	N
2	2	2	2012-06-20	23.00	N

Artefact

ID	Name	Description
1	Vase	Old Vase
2	Knife	Old Knife
3	Pot	Old Pot

Result

ID	Name	Description
3	Pot	Old Pot



MELBOURNE


- List the BuyerID, Name and Phone number for all bidders on artefact 1

```
SELECT * FROM Buyer
WHERE BuyerID IN (SELECT BuyerID FROM Offer
                  WHERE ArtefactID = 1)
```

Equals to

```
SELECT BuyerID, Name and Phone
FROM Buyer NATURAL JOIN Offer
WHERE ArtefactID = 1
```

This is a more efficient way





Exists example

- Returns true if the subquery returns one or more records
- **Example:** *List the BuyerID, Name and Phone number for all bidders on artefact 1*

```
SELECT * FROM Buyer WHERE EXISTS  
(SELECT * FROM Offer WHERE Buyer.BuyerID = Offer.BuyerID  
AND ArtefactID = 1)
```

Offer

SellerID	ArtefactID	BuyerID	Date	Amount	Acceptance
1	1	1	2012-06-20	81223.23	N
1	1	2	2012-06-20	82223.23	N
2	2	1	2012-06-20	19.95	N
2	2	2	2012-06-20	23.00	N

Buyer

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444
3	Oleg	0555555555

Result

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444



- Great tutorial about these:

- <http://www.sqltutorial.org/sql-all/>
- <http://www.sqltutorial.org/sql-any/>
- <http://www.sqltutorial.org/sql-exists/>

- **All:** must satisfy **all** inner conditions

```
SELECT empno, sal
```

```
FROM emp
```

```
WHERE sal > ALL (200, 300, 400);
```

Equiv.

```
SELECT empno, sal
```

```
FROM emp
```

```
WHERE sal > 200 AND sal > 300 AND sal > 400;
```

- **Any:** must satisfy **at least one of** the inner conditions (any of)

```
SELECT empno, sal
```

```
FROM emp
```

```
WHERE sal > ANY (200, 300, 400);
```

Equiv.

```
SELECT empno, sal
```

```
FROM emp
```

```
WHERE sal > 200 OR sal > 300 OR sal > 400;
```

- **Exists:** the inner query returns **at least one** record

```
SELECT empid, first_name, last_name
```

```
FROM employees AS E
```

```
WHERE
```

```
EXISTS( SELECT * FROM dependents AS D WHERE D.empid = E.empid);
```

*“Print all employees who have
at least one dependent”*





- Inserting records from a table:
 - Note: table must already exist

```
INSERT INTO NewEmployee  
SELECT * FROM Employee;
```

- Multiple record inserts:

All columns must be inserted

```
INSERT INTO Employee VALUES  
(DEFAULT, "A", "A's Addr", "2012-02-02", NULL, "S"),  
(DEFAULT, "B", "B's Addr", "2012-02-02", NULL, "S"),  
(DEFAULT, "C", "C's Addr", "2012-02-02", NULL, "S");
```

Specific columns will be inserted

```
INSERT INTO Employee  
(Name, Address, DateHired, EmployeeType)  
VALUES  
("D", "D's Addr", "2012-02-02", "C"),  
("E", "E's Addr", "2012-02-02", "C"),  
("F", "F's Addr", "2012-02-02", "C");
```


- Changes *existing* data in tables
 - Order of statements is important
 - Specifying a WHERE clause is important
 - Unless you want it to operate on the whole table

```
UPDATE Hourly
    SET HourlyRate = HourlyRate * 1.10;
```

- **Example:** *Increase all salaries greater than \$100000 by 10% and all other salaries by 5%*

```
UPDATE Salaried
    SET AnnualSalary = AnnualSalary * 1.05
    WHERE AnnualSalary <= 100000;
UPDATE Salaried
    SET AnnualSalary = AnnualSalary * 1.10
    WHERE AnnualSalary > 100000;
```

Any problems with this?

The UPDATE Statement: CASE

- A better solution in this case is to use the **CASE** command

```
UPDATE Salaried
SET AnnualSalary =
CASE
    WHEN AnnualSalary <= 100000
    THEN AnnualSalary * 1.05
    ELSE AnnualSalary * 1.10
END;
```

If salary is lower than 100000 increase it by 5%,
otherwise increase it by 10%



- REPLACE
 - REPLACE works identically as INSERT
 - Except if an old row in a table has a key value the same as the new row then it is overwritten...
- DELETE
 - The DANGEROUS command – deletes *ALL* records

```
DELETE FROM Employee;
```

 - The better version (unless you are really, really sure)

```
DELETE FROM Employee  
WHERE Name = "Grace";
```
 - Be aware of the foreign key constraints
 - ON DELETE CASCADE or ON DELETE RESTRICT (lab practice)

- Any relation that is not in the physical models, but is made available to the “user” as a virtual relation is called a view.
- Views are good because:
 - They help hide the query complexity from users
 - They help hide data from users
 - Different users use different views
 - Prevents someone from accessing the employee tables to see salaries for instance
 - One way of improving database security
- Create view statement:
CREATE VIEW nameofview **AS** validsqlstatement
- Once a view is defined
 - Its definition is stored in the database (not the data, but metadata – schema information)
 - Can be used just like any other table

Create View Example

MELBOURNE

```
CREATE VIEW EmpPay AS
SELECT Employee.ID, Employee.Name, DateHired,
       EmployeeType, HourlyRate AS Pay
FROM Employee INNER JOIN Hourly
ON Employee.ID = Hourly.ID

UNION

SELECT Employee.ID, Employee.Name, DateHired,
       EmployeeType, AnnualSalary AS Pay
FROM Employee INNER JOIN Salaried
ON Employee.ID = Salaried.ID

UNION

SELECT Employee.ID, Employee.Name, DateHired,
       EmployeeType, BillingRate AS Pay
FROM Employee INNER JOIN Consultant
ON Employee.ID = Consultant.ID;
```



Using a View

SELECT * FROM EmpPay;

ID	Name	DateHired	EmployeeType	Pay
3	Alice	2012-12-02	H	23.43
4	Alan	2010-01-22	H	29.43
1	Sean	2012-02-02	S	92000.00
2	Linda	2011-06-12	S	92300.00
5	Peter	2010-09-07	C	210.00
6	Rich	2012-05-19	C	420.00

**SELECT * FROM EmpPay
WHERE EmployeeType = "H" OR EmployeeType = "C"**

ID	Name	DateHired	EmployeeType	Pay
3	Alice	2012-12-02	H	23.43
4	Alan	2010-01-22	H	29.43
5	Peter	2010-09-07	C	210.00
6	Rich	2012-05-19	C	420.00

Output Snippets Query 1 Result Lecture7.sql Result x

Fetch 4 records. Duration: 0.000 sec, fetched 4 records.



- There are more than CREATE!
- ALTER
 - Allows us to add or remove attributes (columns) from a relation (table)
 - **ALTER TABLE** TableName **ADD** AttributeName AttributeType
 - **ALTER TABLE** TableName **DROP** AttributeName
- RENAME
 - Allows the renaming of tables (relations)
 - **RENAME TABLE** CurrentTableName **TO** NewTableName



MELBOURNE

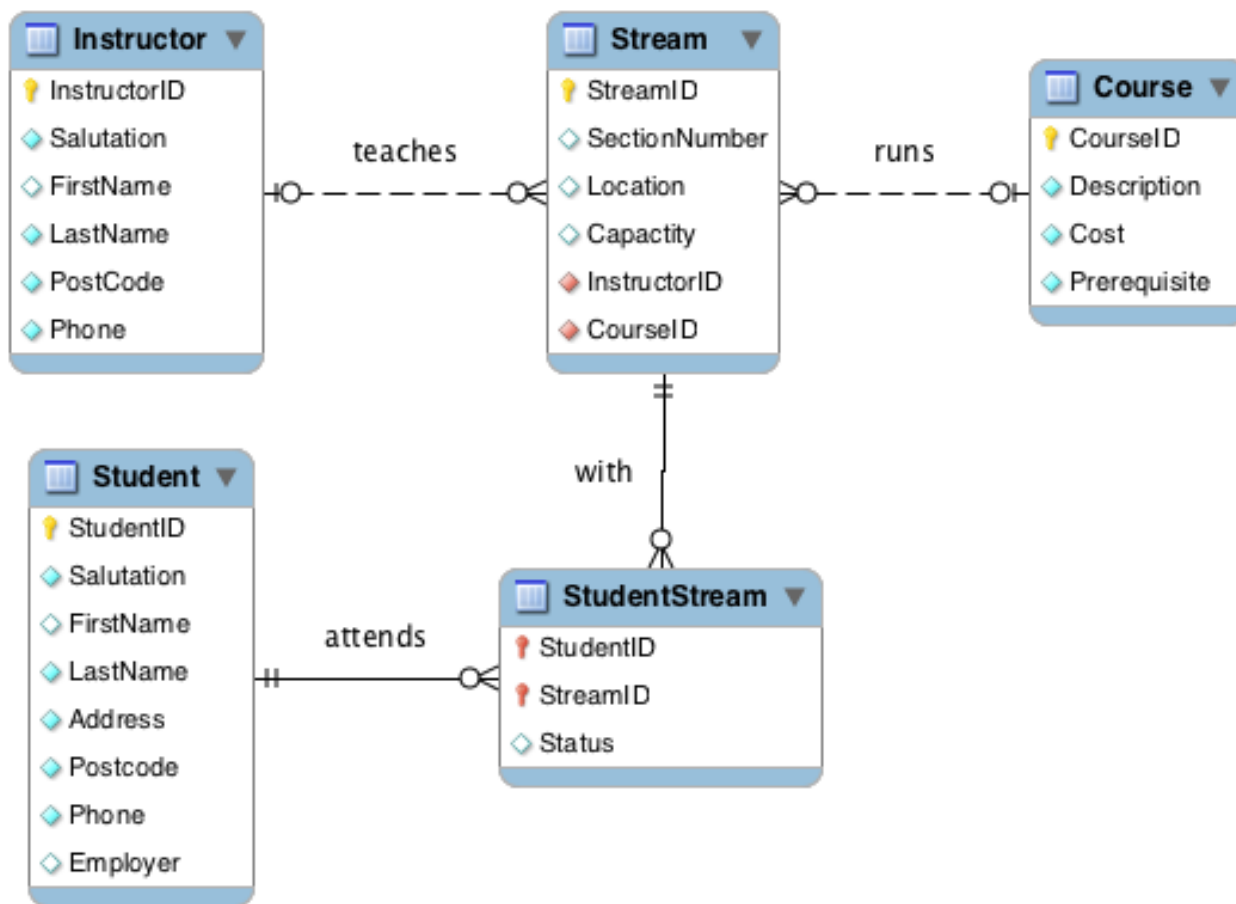
- TRUNCATE
 - Same as DELETE * FROM table;
 - Faster but cannot ROLL BACK a TRUNCATE command
 - Have to get data back from backup...
- DROP
 - Potentially DANGEROUS
 - Kills a relation – removes the data, removes the relation
 - There is NO UNDO COMMAND! (have to restore from backup)
 - DROP TABLE TableName

- DCL
 - Users and permissions
 - **CREATE USER, DROP USER**
 - **GRANT, REVOKE**
 - **SET PASSWORD**
- Other Commands
 - Database administration
 - **BACKUP TABLE, RESTORE TABLE**
 - **ANALYZE TABLE**
 - Miscellaneous
 - **DESCRIBE tablename**
 - **USE db_name**
- They are typically called 'Database Administration Statements'



MELBOURNE

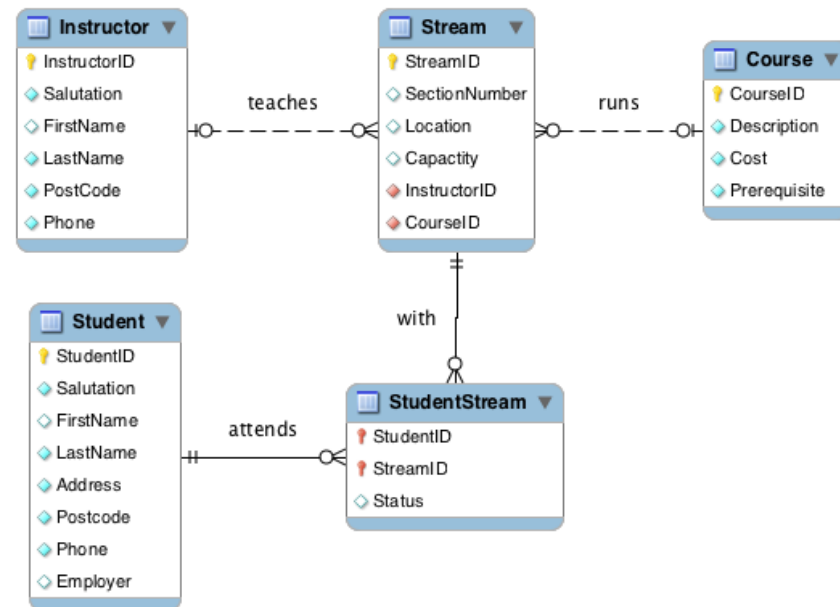
- It's going to be critical for you to think like SQL to handle the queries you will need to write...
- Hopefully the following discussion will help you in this endeavour:
 1. **USE** the database design as a MAP to help you when you are formulating queries
 2. **USE** the structure of the SELECT statement as a template
 3. **FILL** out parts of the SELECT structure and **BUILD** the query
- Let's try it!



Example: Which employers employ students who are doing a course in locations where the capacity is greater than 20 persons, and what are those locations?

Which employers employ students who are doing a course in locations where the capacity is greater than 20 persons, and what are those locations?

- What is the query asking for:
 - Which fields & tables:
F: Employer, Location
T: Student, Stream, StudentStream
 - But only if the capacity > 20 (condition)



- Lets try to use the structure of the SELECT statement now:

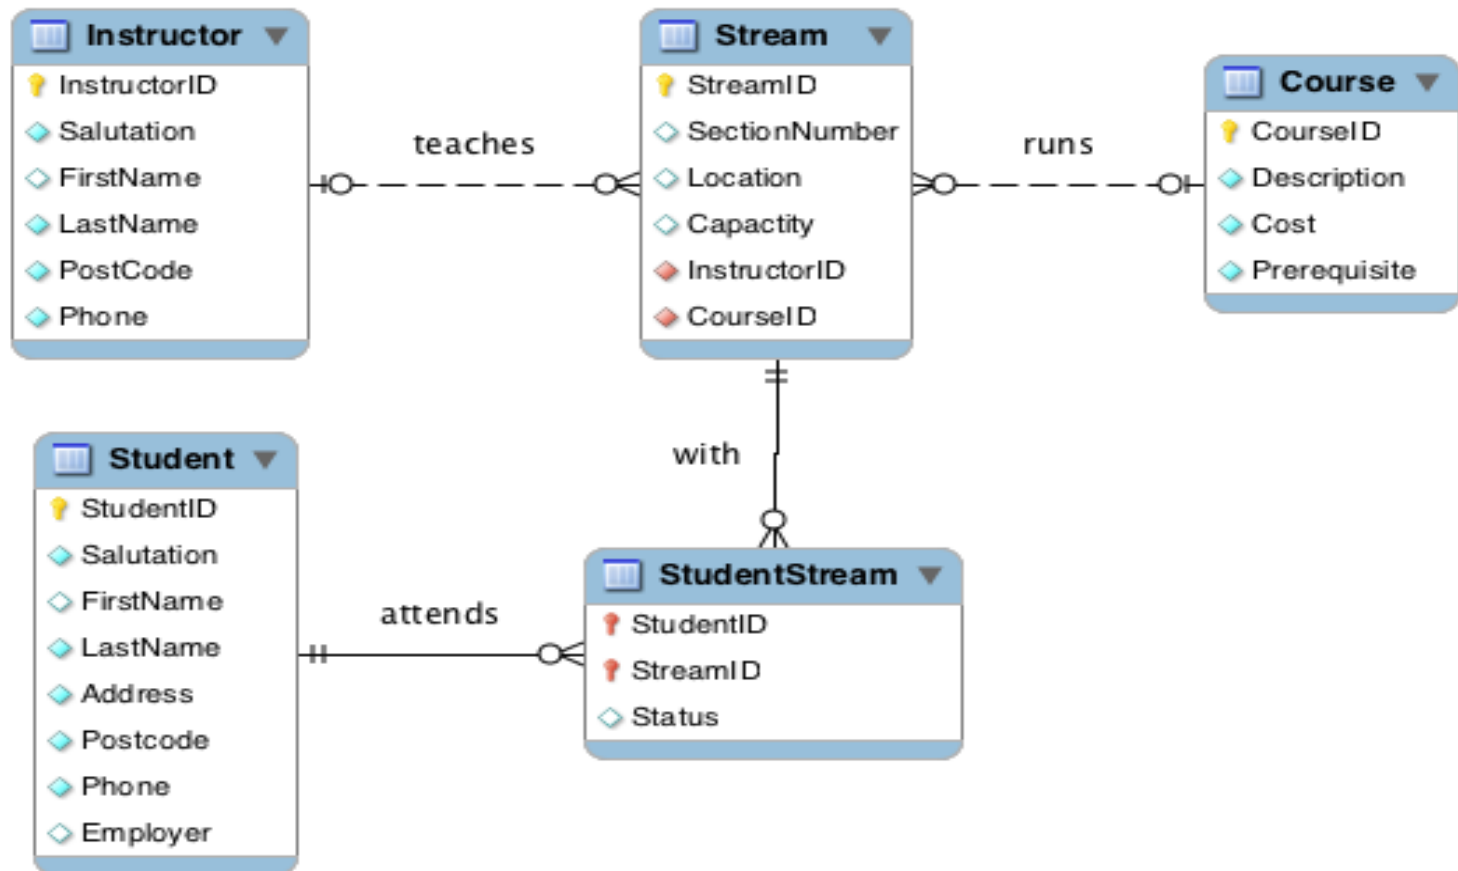
```

SELECT Employer, Location
FROM Student INNER JOIN StudentStream
ON Student.StudentID = StudentStream.StudentID
INNER JOIN Stream
ON StudentStream.StreamID = Stream.StreamID
WHERE Capacity > 20;
  
```

```

SELECT Employer, Location
FROM Student
NATURAL JOIN StudentStream
NATURAL JOIN Stream
WHERE Capacity > 20;
  
```

What is the phone number of the instructor who teaches a course that costs over 10000\$ attended by studentID 202.





MELBOURNE

- A very good overview:

<https://www.youtube.com/watch?v=uRdIdd-UkTc&index=7&list=PLdQddgMBv5zHcEN9RrhADq3CBColhY2hl>



MELBOURNE

- You need to know how to write SQL



MELBOURNE

- Storage and indexing
 - Learn how data is stored and accessed within a DBMS
 - Alternative types of indexes
 - Going “under the hood” of a DBMS